STAT3622 Data Visualization (Lecture 1)

Introduction to Data Science

Dr. Yiwei Fan
The University of Hong Kong

18 January 2021

R Programming

R Programming

- R is a programming language and free software environment for statistical computing and graphics.
- The R language is widely used among statisticians and data miners for developing statistical software and data analysis.
- Although R has a command line interface, there are several third-party graphical user interfaces, such as RStudio, an integrated development environment.

RStudio and R Markdown

RStudio IDE

- RStudio is a popular IDE (Integrated Development Environment) for R programming
- It is a powerful editor for R coding and debugging.
- It is a powerful generator for HTML, PDF, dynamic documents and slide shows.
- RStudio can be run on both Desktop and Cloud.
- Check out more nice features of RStudio at its official website

RStudio IDE

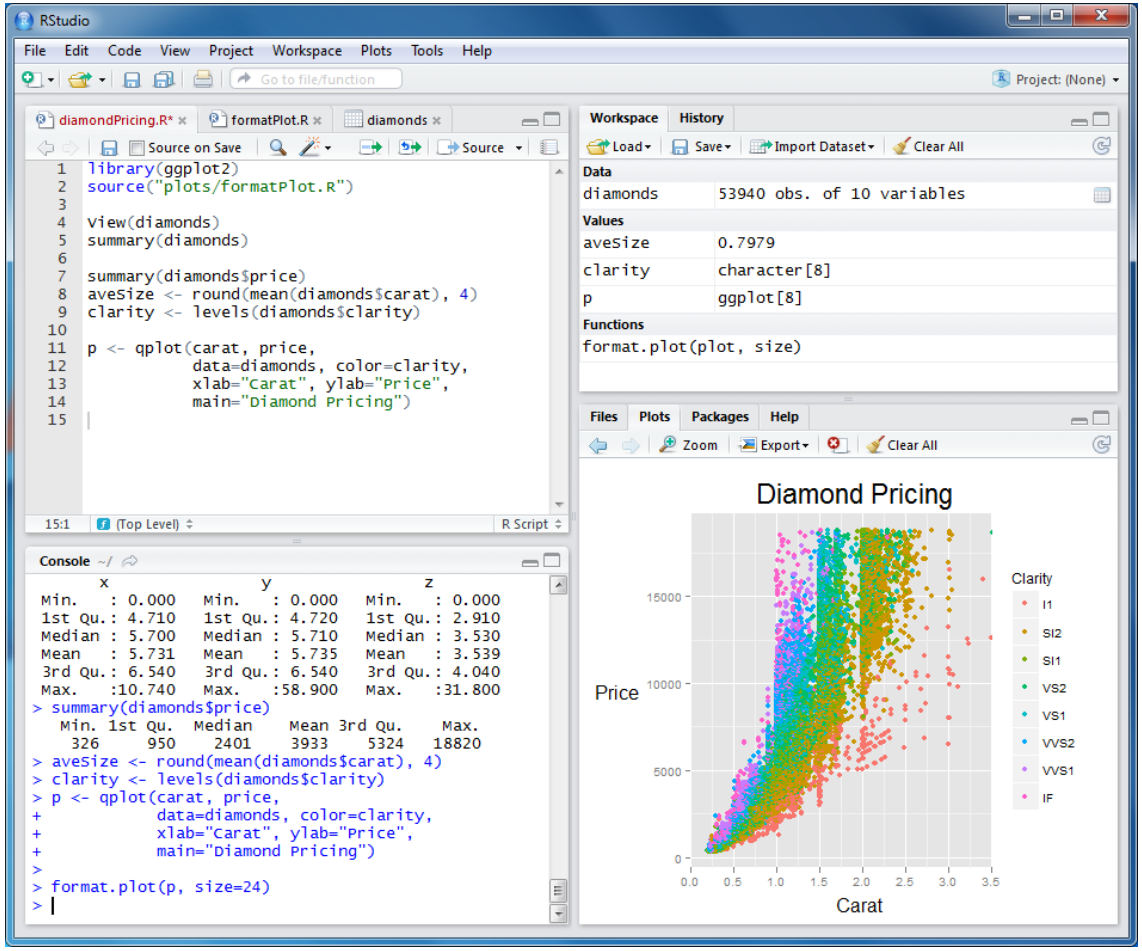
RStudio
File Edit Code View Project Workspace Plots Tools Help
diamondPricing.R* formatPlot.R diamonds
Source on Save Source
1 library(ggplot2)
2 source("plots/formatPlot.R")
3
4 View(diamonds)
5 summary(diamonds)
6
7 summary(diamonds$price)
8 aveSize <- round(mean(diamonds$carat), 4)
9 clarity <- levels(diamonds$clarity)
10
11 p <- qplot(carat, price,
12 data=diamonds, color=clarity,
13 xlab="Carat", ylab="Price",
14 main="Diamond Pricing")
15
15:1 (Top Level) R Script
Console ~/
x y z
Min. : 0.000 Min. : 0.000 Min. : 0.000
1st Qu.: 4.710 1st Qu.: 4.720 1st Qu.: 2.910
Median : 5.700 Median : 5.710 Median : 3.530
Mean : 5.731 Mean : 5.735 Mean : 3.539
3rd Qu.: 6.540 3rd Qu.: 6.540 3rd Qu.: 4.040
Max. :10.740 Max. :58.900 Max. :31.800
> summary(diamonds$price)
Min. 1st Qu. Median Mean 3rd Qu. Max.
326 950 2401 3933 5324 18820
> aveSize <- round(mean(diamonds$carat), 4)
> clarity <- levels(diamonds$clarity)
> p <- qplot(carat, price,
+ data=diamonds, color=clarity,
+ xlab="Carat", ylab="Price",
+ main="Diamond Pricing")
>
> format.plot(p, size=24)
>
Workspace History
Load Save Import Dataset Clear All
Data
diamonds 53940 obs. of 10 variables
Values
aveSize 0.7979
clarity character[8]
p ggplot[8]
Functions
format.plot(plot, size)
Files Plots Packages Help
Zoom Export Clear All
Diamond Pricing
Price
15000
10000
5000
0
0.0 0.5 1.0 1.5 2.0 2.5 3.0 3.5
Carat
Clarity
I1
SI2
SI1
VS2
VS1
VVS2
VVS1
IF

R Markdown

- Click here to view a fantastic micro-video tutorial
- Browse here for a gallery of creative Rmarkdown works

R Markdown (Demonstrated)

```
knitr::kable(head(iris), format = 'html')
```

Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
5.1	3.5	1.4	0.2	setosa
4.9	3.0	1.4	0.2	setosa
4.7	3.2	1.3	0.2	setosa
4.6	3.1	1.5	0.2	setosa
5.0	3.6	1.4	0.2	setosa
5.4	3.9	1.7	0.4	setosa

- Dynamic documentation: report, table, graphics ...
- R packages by Yihui Xie: knitr, bookdown, xaringan, etc

R Markdown (Demonstrated)

```
plot(iris, col=iris$Species)
```

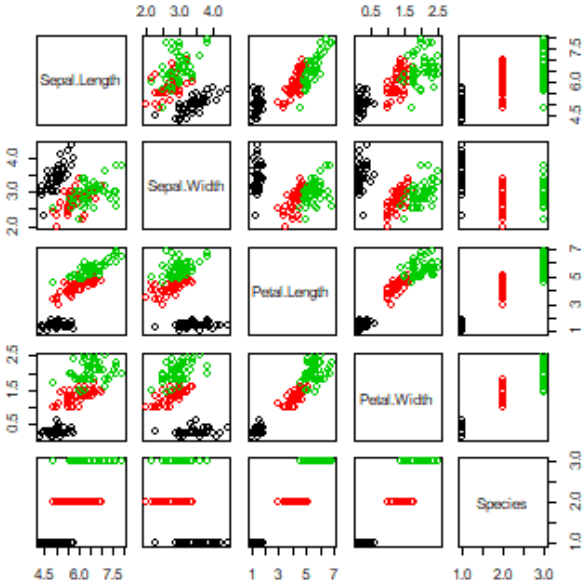

- Data-generated graphics that are reproducible

Exploratory Data Analysis

Exploratory Data Analysis

The EDA is a statistical approach to make sense of data by using a variety of techniques (mostly graphical). It may help us

- Assess assumption about variables distribution
- Identify relationship between variables
- Extract important variables
- Suggest use of appropriate models
- Detect problems of collected data (e.g. outliers, missing data, measurement errors)

Example: Anscombe Dataset

Anscombe Dataset:

x1	y1	x2	y2	x3	y3	x4	y4
10	8.04	10	9.14	10	7.46	8	6.58
8	6.95	8	8.14	8	6.77	8	5.76
13	7.58	13	8.74	13	12.74	8	7.71
9	8.81	9	8.77	9	7.11	8	8.84
11	8.33	11	9.26	11	7.81	8	8.47
14	9.96	14	8.10	14	8.84	8	7.04
6	7.24	6	6.13	6	6.08	8	5.25
4	4.26	4	3.10	4	5.39	19	12.50
12	10.84	12	9.13	12	8.15	8	5.56
7	4.82	7	7.26	7	6.42	8	7.91
5	5.68	5	4.74	5	5.73	8	6.89

Source: Anscombe, F. J. (1973). Graphs in statistical analysis. *American Statistician*, **27**, 17-21.

Example: Anscombe Dataset (Descriptive)

Mean and standard deviation:

	x1	y1	x2	y2	x3	y3	x4	y4
mean	9.00	7.50	9.00	7.50	9.00	7.50	9.00	7.50
sd	3.32	2.03	3.32	2.03	3.32	2.03	3.32	2.03

x-y correlation:

rho1	rho2	rho3	rho4
0.82	0.82	0.82	0.82

Example: Anscombe Dataset (Graphic)

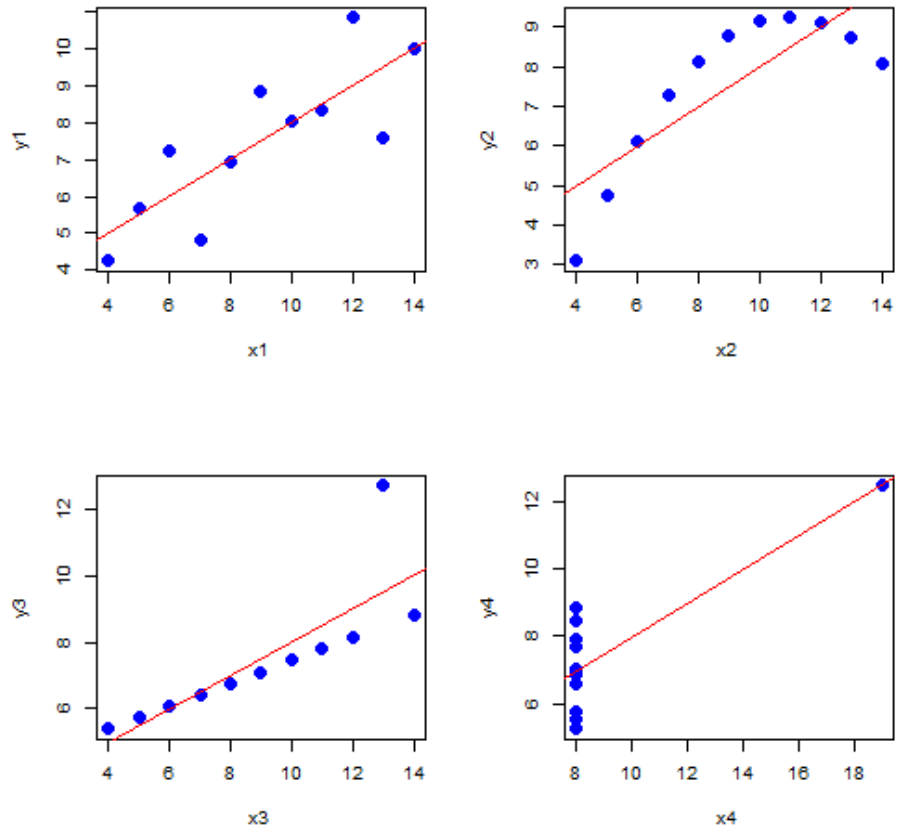

Statistical Graphics

- **Univarite**
  - Histogram, Stem-and-Leaf, Dot, Q-Q, Density plots
  - Boxplot, Box-and-whisker
  - Bar, Pie, Polar, Waterfall charts
- **Bivariate**
  - XYplot, Line, Area, Scatter, Bubble charts
- **Trivariate**
  - 3D Scatter, Contour, Level/Heatmap, Surface plots

Which Chart to Use?

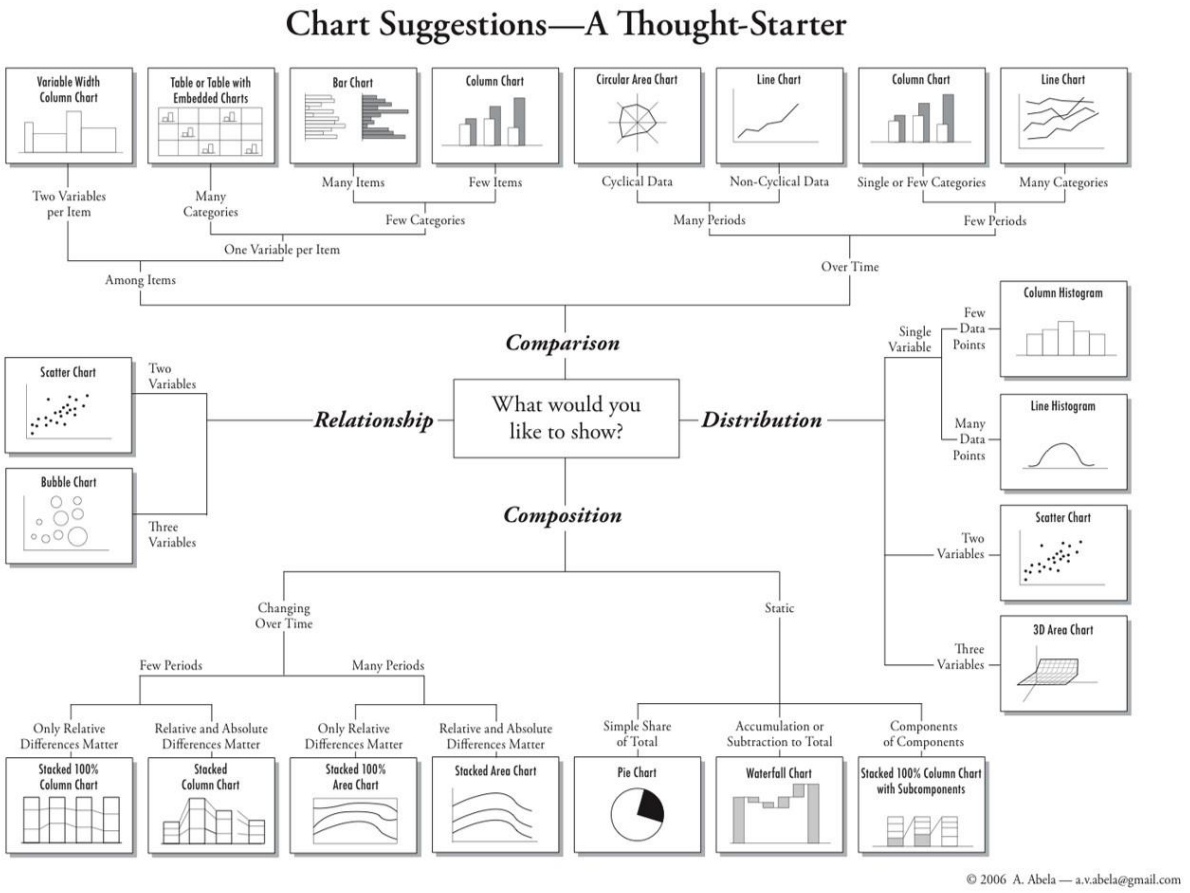

Simple Base Graphics

Iris Dataset

iris setosa	iris versicolor	iris virginica

```
DataX = iris   # ?iris
str(DataX)
```

```
## 'data.frame':    150 obs. of  5 variables:
##  $ Sepal.Length: num  5.1 4.9 4.7 4.6 5 5.4 4.6 5 4.4 4.9 ...
##  $ Sepal.Width : num  3.5 3 3.2 3.1 3.6 3.9 3.4 3.4 2.9 3.1 ...
##  $ Petal.Length: num  1.4 1.4 1.3 1.5 1.4 1.7 1.4 1.5 1.4 1.5 ...
##  $ Petal.Width : num  0.2 0.2 0.2 0.2 0.2 0.4 0.3 0.2 0.2 0.1 ...
##  $ Species     : Factor w/ 3 levels "setosa","versicolor",..: 1 1 1 1 1 1 1 1 1 1 ...
```

```
dim(DataX)
```

```
## [1] 150   5
```

```
head(DataX)  # tail
```

```
##   Sepal.Length Sepal.Width Petal.Length Petal.Width Species
## 1          5.1         3.5          1.4         0.2  setosa
## 2          4.9         3.0          1.4         0.2  setosa
## 3          4.7         3.2          1.3         0.2  setosa
## 4          4.6         3.1          1.5         0.2  setosa
## 5          5.0         3.6          1.4         0.2  setosa
## 6          5.4         3.9          1.7         0.4  setosa
```

```
summary(DataX)
```

```
##   Sepal.Length    Sepal.Width     Petal.Length    Petal.Width          Species
##  Min.   :4.300   Min.   :2.000   Min.   :1.000   Min.   :0.100   setosa    :50
##  1st Qu.:5.100   1st Qu.:2.800   1st Qu.:1.600   1st Qu.:0.300   versicolor:50
##  Median :5.800   Median :3.000   Median :4.350   Median :1.300   virginica :50
##  Mean   :5.843   Mean   :3.057   Mean   :3.758   Mean   :1.199
##  3rd Qu.:6.400   3rd Qu.:3.300   3rd Qu.:5.100   3rd Qu.:1.800
##  Max.   :7.900   Max.   :4.400   Max.   :6.900   Max.   :2.500
```

Basic R Plots: Histogram and Density Plot

```
x = DataX$Sepal.Length # a continuous variable
par(mfrow=c(1,3))
hist(x, main='Histogram (Default)')
hist(x, breaks=20, col=5, main='More Bins and Coloring')
hist(x, breaks=20, freq=F, main='Histogram plus Density Plot')   # using freq=FALSE
lines(density(x), col=2, lty=1, lwd=2)  #add the density curve
```

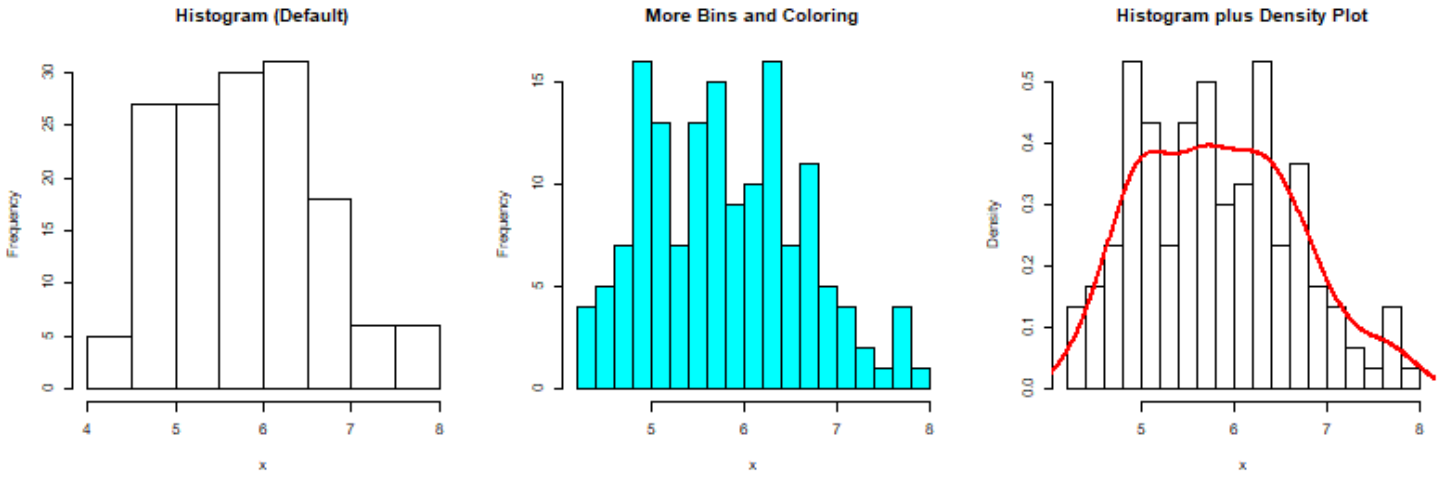

Basic R Plots: Boxplot

```
par(mfrow=c(1,3))
boxplot(DataX$Sepal.Width, main='Boxplot of Sepal.Width')  # Outliers
boxplot(DataX[,1:4], col=c(2,3,4,5), main='Side-by-side Boxplot')
boxplot(Sepal.Width~Species, DataX, col=c(6,7,8), main="Boxplot with Grouping")
```

- An outlier is an observation that is numerically distant from the rest of the data (e.g: outside 1.5 times the interquartile range above the upper quartile and below the lower quartile).

Basic R Plots: Pie and Bar Charts

```
DataX$Flag = DataX$Sepal.Length>5 # Create a binary flag
par(mfrow=c(1,3))
pie(table(DataX$Species[DataX$Flag]), col=c(2,3,4))
barplot(table(DataX$Species[DataX$Flag]), col=c(5,6,7))
barplot(table(DataX$Species, DataX$Flag), col=c(5,6,7), beside=T)
```

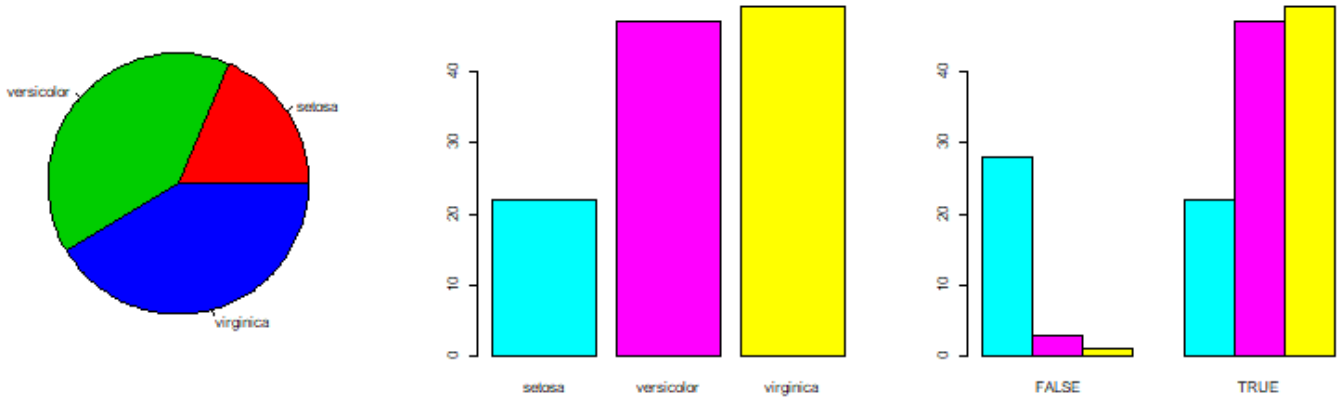

Relationship Between Variables

```
x = DataX$Petal.Length; y = DataX$Petal.Width; z = DataX$Species
par(mfrow=c(1,2)); par(mar=c(4,4,1,4))
plot(x, y, xlab="Petal.Length", ylab="Petal.Width")
abline(coef(lm(y~x)), col=1, lty=2)
plot(x, y, col=c(2,3,4)[z], pch=20, cex=2.0, xlab="Petal.Length", ylab="Petal.Width")
abline(lm(y~x), col=1, lty=2)
legend("topleft", levels(z), pch=20, col=c(2,3,4))
```

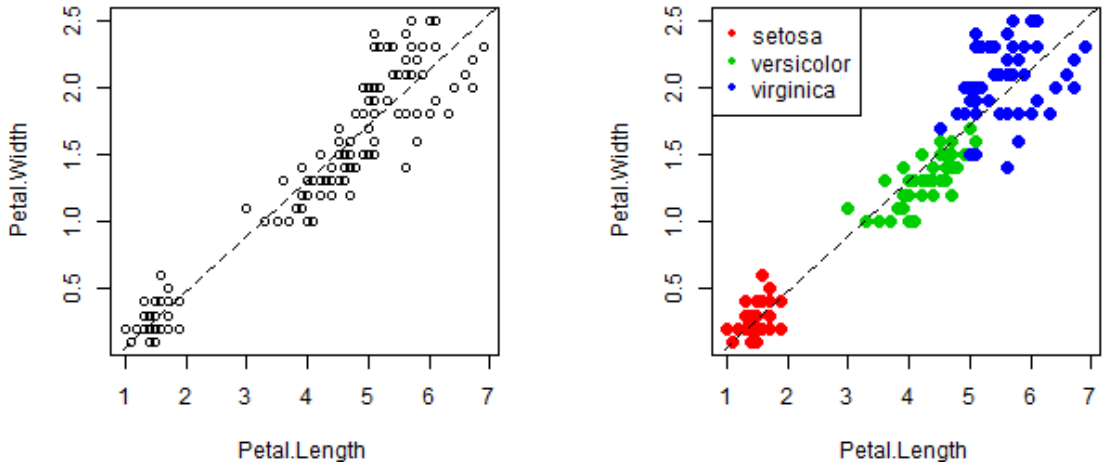

Pairwise Scatter Plot

```
plot(DataX, col=DataX$Species,
     main="Pairwise Scatter Plot")
```

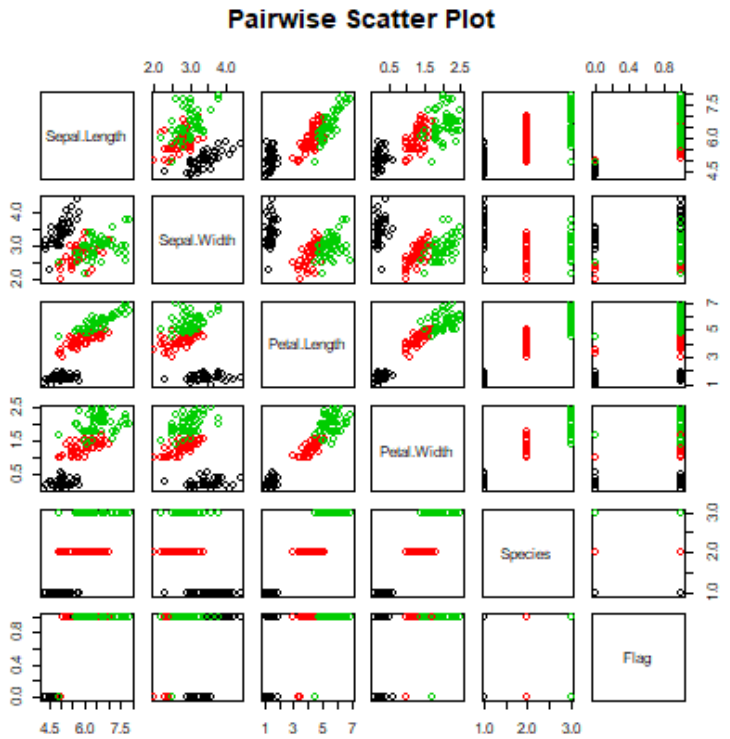


```
pairs(DataX[,1:4],
      col = c(4,5,6)[DataX$Species],
      main="More Sophisticated")
```

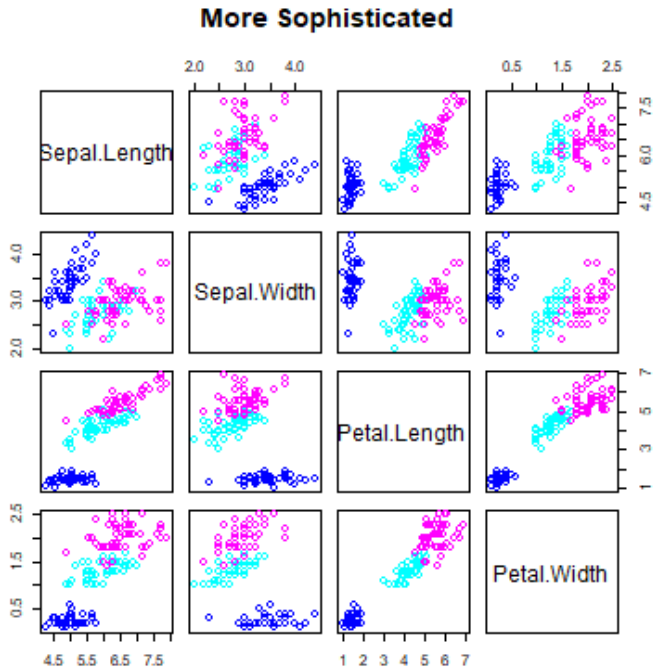

Using R:Lattice Package

R:Lattice

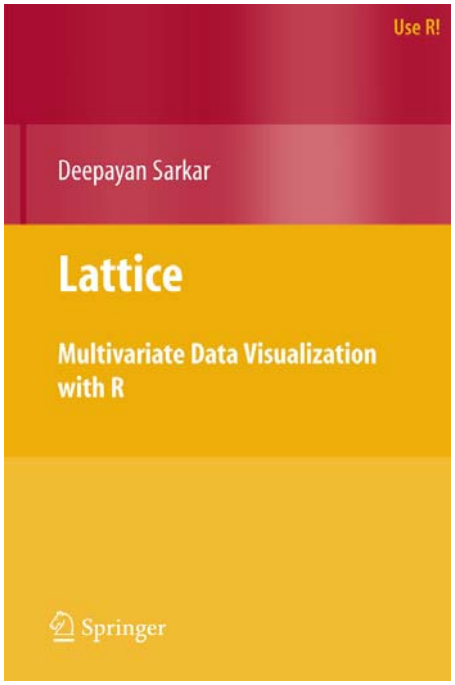

- Using trellis graphs for multivariate data
- Multipanel conditioning and grouping
- Elegant high-level data visualization
- Covering most of statistical charts
- Figures and Codes can be found at http://lmdvr.r-forge.r-project.org/
- However, plot customization are not so straightforward

Univariate Distributions

```
library(lattice); library(gridExtra)
p1 = histogram(DataX$Sepal.Length)
p2 = bwplot(DataX$Sepal.Length)
grid.arrange(p1, p2, ncol=2)
```

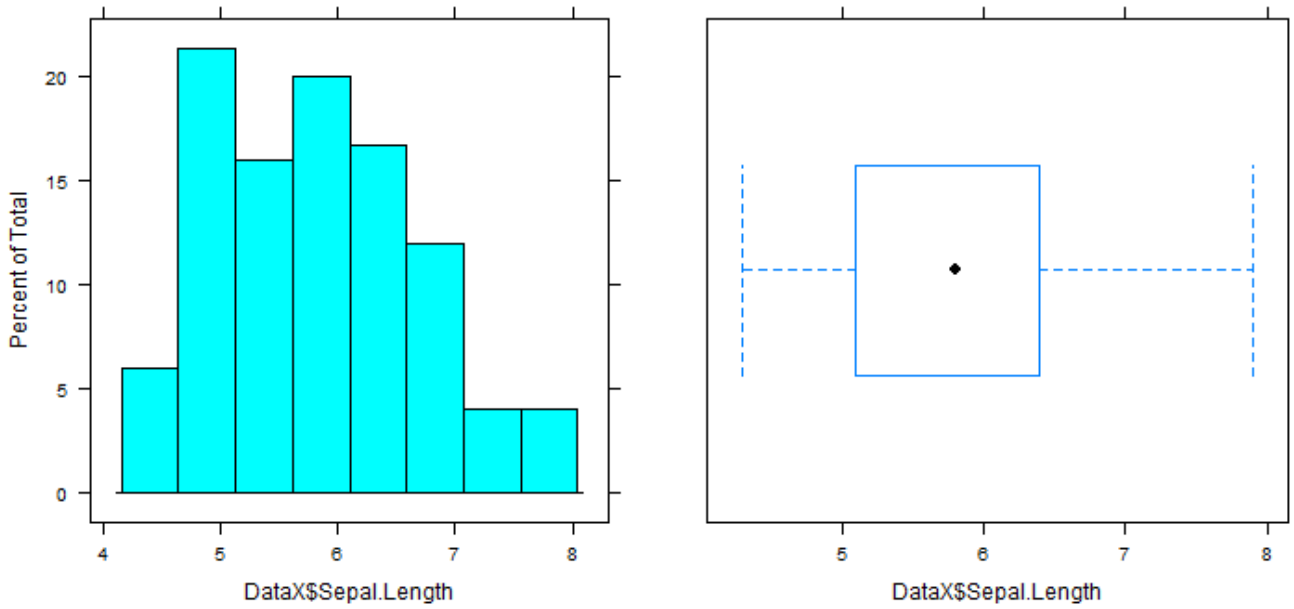

Histogram with Conditioning

```
histogram(data=DataX, ~Sepal.Length|Species, breaks=12, layout = c(3, 1))
```

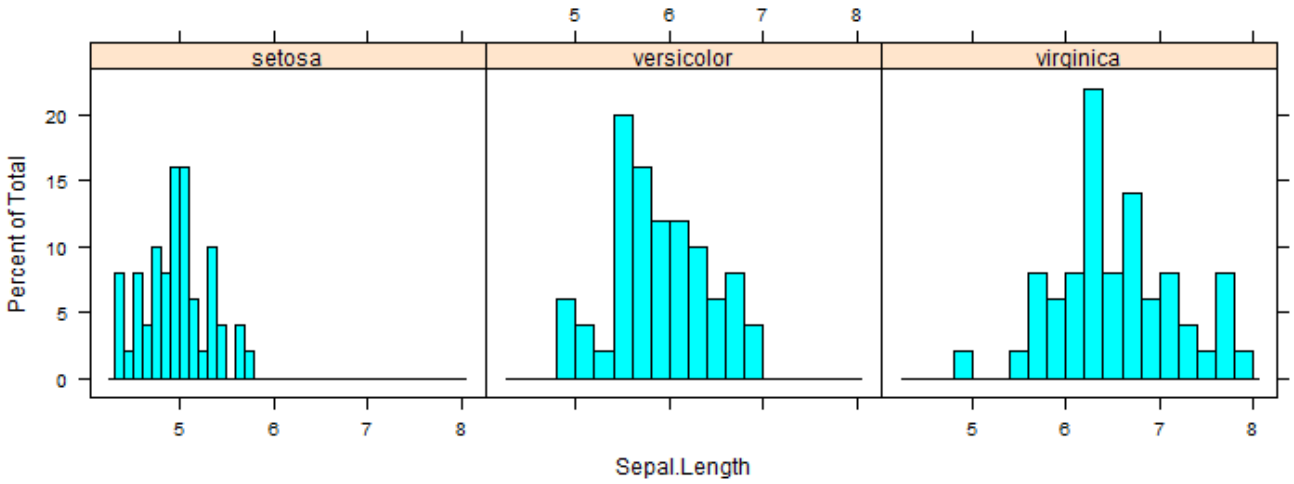

Density plot with Grouping

```
densityplot(data=DataX, ~Sepal.Length, groups=Species,
            plot.points=F, auto.key=list(space="top", columns=3))
```

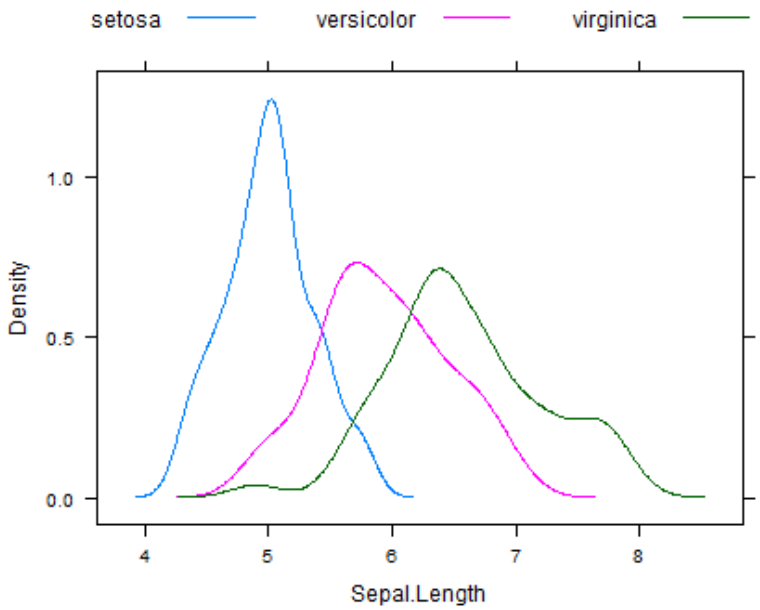

Boxplot with Grouping

```
bwplot(data=DataX, Sepal.Width~Species)
```

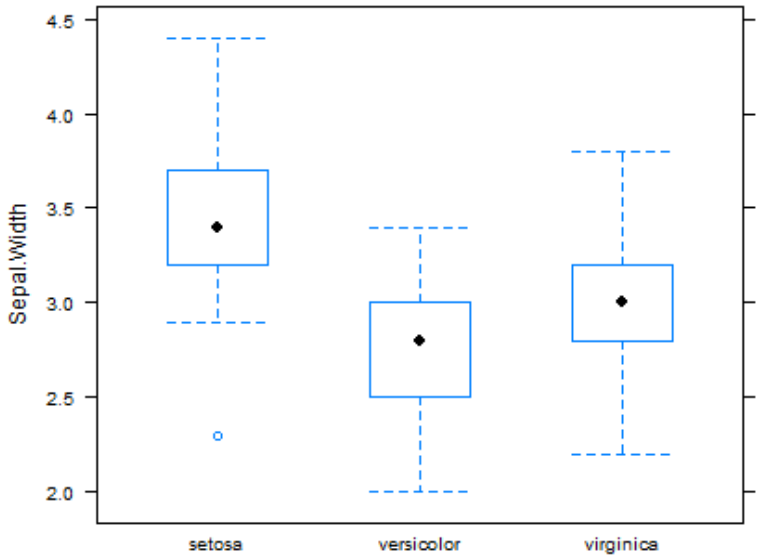

Bivariate plot with Grouping

```
xyplot(data=DataX, Sepal.Length ~ Petal.Length, groups = Species,
       type = c("p", "smooth",  "g"),
       auto.key = list(space="top", columns=3)) # grouping
```

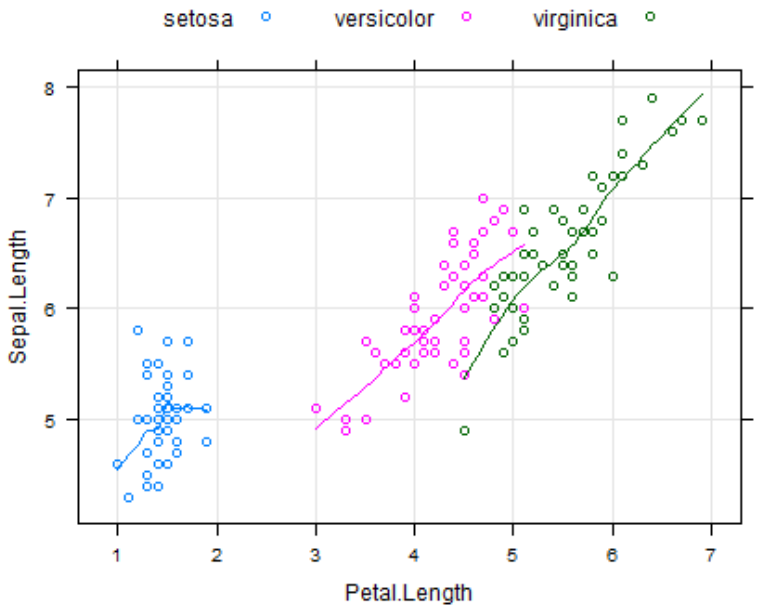

Bivariate plot with Conditioning

```
xyplot(data=DataX, Sepal.Length ~ Petal.Length|Species,
       type=c("p", "smooth",  "g"), layout=c(1,3)) # conditioning
```

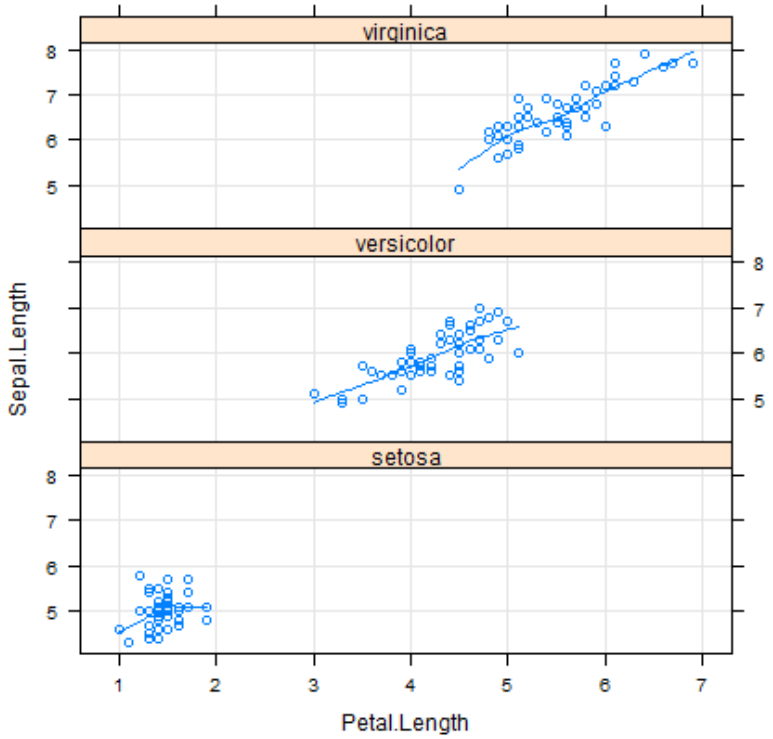

Trivariate 3D Plot

```
cloud(data=DataX, Sepal.Length ~ Sepal.Length * Petal.Width, groups = Species,
      auto.key = list(space="top", columns=3), panel.aspect = 0.8)
```

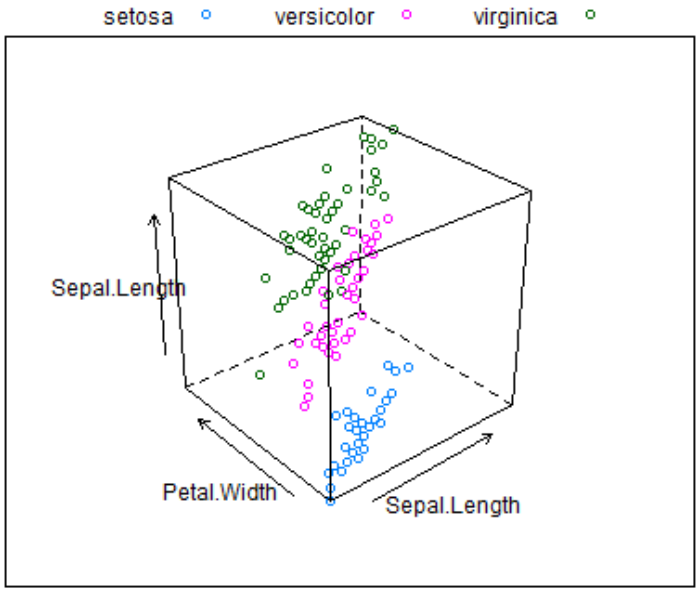

Trivariate Heatmap

```
dist = as.matrix(dist(DataX[,3:4]))
levelplot(dist, colorkey = T, col.regions = terrain.colors,
          scales = list(at=c(0,0),tck = c(0,0)),
          xlab="",ylab="",main="Levelplot of Pairwise Distance Matrix")
```

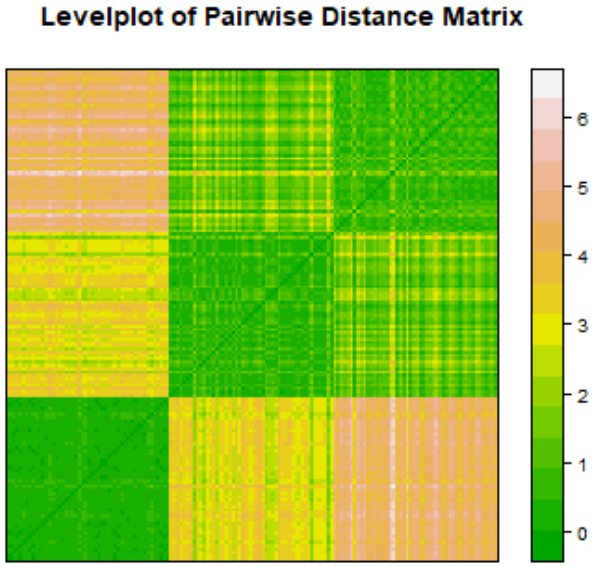

Thank you!